



PURDUE UNIVERSITY NORTHWEST

PROJECT

DEEP LEARNING CLASS

PROFESSOR RICARDO CALIX

***APPLYING CONVOLUTION NEURAL NETWORK ON COLOUR
IMAGES DATASET.***

DATA SOURCE: KAGGLE DATASETS

DATA SET:

CLASSIFICATION OF SCENERY IMAGES:

Objective:

***To classify scenery images that consist of 4 classes that are sea ,
mountain , forest and buildings.***

SAMPLE IMAGES :



CONVOLUTION NEURAL NETWORK CODE:

```
#!/usr/bin/env python
```

```
## Deep Learning code
```

```
##Professor Ricardo Calix
```

```
## simple example of a convolutional neural network for RGB data
```

```
#####
```

```
import sklearn
```

```
import tensorflow as tf
```

```

import numpy as np

from numpy import genfromtxt

from sklearn import datasets

#from sklearn.cross_validation import train_test_split

from sklearn.model_selection import train_test_split

from sklearn.metrics import accuracy_score

from sklearn.metrics import confusion_matrix

from sklearn.metrics import precision_score

from sklearn.metrics import recall_score, f1_score

import pandas as pd

import matplotlib.pyplot as plt

#####

## set parameters


import warnings

warnings.filterwarnings("ignore")

np.set_printoptions(threshold=np.inf) #print all values in numpy array

#####

#parameters for the main loop


learning_rate = 0.001

n_epochs = 6000 ##27000

batch_size = 150

display_step = 10


# Parameters for the network

```

```
n_input = 7500 # 50x50x3
```

```
n_classes = 4 # (0-3 classes)
```

```
dropout = 0.75 # Dropout, probability to keep units
```

```
#####  
#####
```

```
## this will create your own data set (i.e. your_mnist)
```

```
## put your images in testA (I used pngs)
```

```
from PIL import Image
```

```
#import cv2
```

```
import glob
```

```
import numpy as np
```

```
train = []
```

```
labels = []
```

```
#####  
#####
```

```
files = glob.glob ("scenery/forest/*.jpg") # your images path
```

```
for myFile in files:
```

```
    my_im = Image.open(myFile).convert('RGB')  ## .convert('LA') ## is for greyscale
```

```
    #my_im.show()
```

```
    resized_my_im = my_im.resize((50,50))  ## resize from 150x150x3 to 50x50x3
```

```
#resized_my_im.show()
```

```
image = np.array(resized_my_im)
```

```
#image = np.array(my_im)
```

```
print(image.shape)
```

```
new_image = image.reshape(image.shape[0]*image.shape[1]*image.shape[2])
```

```
print(new_image.shape)
```

```
#input_string = input("??")
```

```
train.append(new_image)
```

```
labels.append(0)
```

```
#####  
#####
```

```
files = glob.glob ("scenery/buildings/*.jpg") # your images path
```

```
for myFile in files:
```

```
    my_im = Image.open(myFile).convert('RGB')  ## .convert('LA') ## is for greyscale
```

```
    #my_im.show()
```

```
    resized_my_im = my_im.resize((50,50))  ## resize from 150x150x3 to 50x50x3
```

```
    #resized_my_im.show()
```

```
    image = np.array(resized_my_im)
```

```
    #image = np.array(my_im)
```

```
    print(image.shape)
```

```
    new_image = image.reshape(image.shape[0]*image.shape[1]*image.shape[2])
```

```
    print(new_image.shape)
```

```
    #input_string = input("??")
```

```
    train.append(new_image)
```

```
    labels.append(1)
```

```
#####  
#####
```

```
files = glob.glob ("scenery/mountain/*.jpg") # your images path
```

```
for myFile in files:
```

```
    my_im = Image.open(myFile).convert('RGB')  ## .convert('LA') ## is for greyscale
```

```
    #my_im.show()
```

```
    resized_my_im = my_im.resize((50,50))  ## resize from 150x150x3 to 50x50x3
```

```
    #resized_my_im.show()
```

```
    image = np.array(resized_my_im)
```

```
    #image = np.array(my_im)
```

```
    print(image.shape)
```

```
    new_image = image.reshape(image.shape[0]*image.shape[1]*image.shape[2])
```

```
    print(new_image.shape)
```

```
    #input_string = input("??")
```

```
    train.append(new_image)
```

```
    labels.append(2)
```

```
#####  
#####
```

```
files = glob.glob ("scenery/sea/*.jpg") # your images path
```

```
for myFile in files:
```

```
    my_im = Image.open(myFile).convert('RGB')  ## .convert('LA') ## is for greyscale
```

```
    #my_im.show()
```

```

resized_my_im = my_im.resize((50,50))  ## resize from 150x150x3 to 50x50x3

#resized_my_im.show()

image = np.array(resized_my_im)

#image = np.array(my_im)

print(image.shape)

new_image = image.reshape(image.shape[0]*image.shape[1]*image.shape[2])

print(new_image.shape)

#input_string = input("??")

train.append(new_image)

labels.append(3)

#####

#####

#####

#####

train = np.array(train,dtype='float32')

labels = np.array(labels, dtype='float32')

# convert (number of images x height x width x number of channels) to (number of images x (height *
width *3))

# for example (120 * 40 * 40 * 3)-> (120 * 4800)

#train = np.reshape(train,[train.shape[0],train.shape[1]*train.shape[2]*train.shape[3]])

#train = np.reshape(train,[train.shape[0],train.shape[1]*train.shape[2]])

print(train.shape)

print(labels.shape)

#print(train.shape[1])

#input_string = input("train size is")

```

```
#####  
#####
```

```
# save numpy array as .npy formats
```

```
np.save('train',train)
```

```
your_mnist = train
```

```
#####  
#####
```

```
## normalization is very important
```

```
x=your_mnist
```

```
xmax, xmin = x.max(), x.min()
```

```
x = (x - xmin)/(xmax - xmin)
```

```
your_mnist = x
```

```
#####  
#####
```

```
x_all = your_mnist
```

```
labels_all = labels
```

```
x_train, x_test, y_train, y_test = train_test_split(x_all, labels_all, test_size=0.30, random_state=42)
```



```
#####  
#####
```

```
## print stats
```

```
precision_scores_list = []
```

```
accuracy_scores_list = []
```

```
def print_stats_metrics(y_test, y_pred):
```

```
    print('Accuracy: %.2f' % accuracy_score(y_test, y_pred) )
```

```
    #Accuracy: 0.84
```

```
    accuracy_scores_list.append(accuracy_score(y_test, y_pred) )
```

```
    confmat = confusion_matrix(y_true=y_test, y_pred=y_pred)
```

```
    print "confusion matrix"
```

```
    print(confmat)
```

```
    print pd.crosstab(y_test, y_pred, rownames=['True'], colnames=['Predicted'], margins=True)
```

```
    precision_scores_list.append(precision_score(y_true=y_test, y_pred=y_pred, average='weighted'))
```

```
    print('Precision: %.3f' % precision_score(y_true=y_test, y_pred=y_pred, average='weighted'))
```

```
    print('Recall: %.3f' % recall_score(y_true=y_test, y_pred=y_pred, average='weighted'))
```

```
    print('F1-measure: %.3f' % f1_score(y_true=y_test, y_pred=y_pred, average='weighted'))
```

```
#####
```

```
#####
```

```
def conv2d(x, W, b, strides=1):
```

```
    # Conv2D function, with bias and relu activation
```

```
    x = tf.nn.conv2d(x, W, strides=[1, strides, strides, 1], padding='SAME')
```

```
    x = tf.nn.bias_add(x, b)
```

```
    return tf.nn.relu(x)    ## relu removes negative values
```

```
#####
```

```
def maxpool2d(x, k=2):
```

```
    # MaxPool2D function
```

```
    # padding='SAME' is very useful for uneven images. If maxpooling
```

```
    # image 25x25 -> 12.5 x 12.5 then it is rounded up to 13x13
```

```
    return tf.nn.max_pool(x, ksize=[1, k, k, 1], strides=[1, k, k, 1],
```

```
        padding='SAME')
```

```
#####
```

```
def layer(input, weight_shape, bias_shape):
```

```
    W = tf.Variable(tf.random_normal(weight_shape))
```

```
    b = tf.Variable(tf.random_normal(bias_shape))
```

```
    mapping = tf.matmul(input, W)
```

```
    result = tf.add( mapping , b )
```

```
    return result
```

```
#####
```

```
def conv_layer(input, weight_shape, bias_shape):
```

```
    ##rr =raw_input()
```

```
    W = tf.Variable(tf.random_normal(weight_shape))
```

```
    b = tf.Variable(tf.random_normal(bias_shape))
```

```
    conv = conv2d(input, W, b)
```

```
    # Max Pooling (down-sampling)
```

```
conv_max = maxpool2d(conv, k=2)
```

```
return conv_max
```

```
#####
```

```
def fully_connected_layer(conv_input, fc_weight_shape, fc_bias_shape, dropout):
```

```
    new_shape = [-1, tf.Variable(tf.random_normal(fc_weight_shape)).get_shape().as_list()[0]]
```

```
    fc = tf.reshape(conv_input, new_shape)
```

```
    w_fc = tf.Variable( tf.random_normal( fc_weight_shape ) )
```

```
    mapping = tf.matmul( fc , w_fc ) # y = w * x
```

```
    fc = tf.add( mapping, tf.Variable(tf.random_normal( fc_bias_shape )) )
```

```
    fc = tf.nn.relu(fc)
```

```
    # Apply Dropout
```

```
    fc = tf.nn.dropout(fc, dropout)
```

```
    return fc
```

```
#####
```

```
## define the architecture here
```

```
def inference_conv_net_3_convolutions(x, dropout):
```

```
    # Reshape input picture
```

```
    # shape = [-1, size_image_x, size_image_y, 3 channels (e.g. rgb)]
```

```
    # the image for rgb and batches of 150 would be [150, 7500] because
```

```
    # there are 128 samples per batch and images are 50x50x3 = 7500
```

```
    # this has to be re-shaped because Convolutional layers only take 4 dimensional tensors as input
```

```
    # the -1 infers the number of batches and then we make the 7500 into 50x50x3
```

```
    x = tf.reshape(x, shape=[-1, 50, 50, 3])
```

```
# Convolution Layer 1, filter 5x5 conv, 3 inputs or 3 channels, 16 outputs
```

```
# max pool will reduce image from 50x50 to 25x25
```

```
conv1 = conv_layer(x, [5, 5, 3, 16], [16] )
```

```
# Convolution Layer 2, filter 5x5 conv, 16 inputs, 36 outputs
```

```
# max pool will reduce image from 25x25 to 13x13
```

```
conv2 = conv_layer(conv1, [5, 5, 16, 36], [36] )
```

```
# Convolution Layer 2, filter 5x5 conv, 36 inputs, 64 outputs
```

```
# max pool will reduce image from 13x13 to 7x7
```

```
conv3 = conv_layer(conv2, [5, 5, 36, 64], [64] )
```

```
# Fully connected layer, 7*7*64 inputs, 1024 outputs
```

```
# Reshape conv2 output to fit fully connected layer input
```

```
# maxpool function padding=same rounds up 6.5 to 7
```

```
fc1 = fully_connected_layer(conv3, [7*7*64, 1024], [1024] , dropout)
```

```
# Output, 128 inputs, 10 outputs (class prediction)
```

```
output = layer(fc1 ,[1024, n_classes], [n_classes] )
```

```
return output
```

```
#####
```

```
def loss_deep_conv_net(output, y_tf):
```

```
    xentropy = tf.nn.softmax_cross_entropy_with_logits(logits=output, labels=y_tf)
```

```
    loss = tf.reduce_mean(xentropy)
```

```
    return loss
```

```
#####
```

```
def training(cost):
```

```
    optimizer = tf.train.AdamOptimizer(learning_rate=learning_rate)
```

```
    train_op = optimizer.minimize(cost)
```

```
    return train_op
```

```
#####
```

```
def evaluate(output, y_tf):
```

```
    correct_prediction = tf.equal(tf.argmax(output,1), tf.argmax(y_tf,1))
```

```
    accuracy = tf.reduce_mean(tf.cast(correct_prediction, "float"))
```

```
    return accuracy
```

```
#####
```

```
x_tf = tf.placeholder(tf.float32, [None, n_input]) ## 50x50x3
```

```
y_tf = tf.placeholder(tf.float32, [None, n_classes])
```

```
keep_prob = tf.placeholder(tf.float32) #dropout (keep probability)
```

```
#####
```

```
output = inference_conv_net_3_convolutions(x_tf, keep_prob)
```

```
#output = inference_conv_net2(x_tf, keep_prob)
```

```
#output = inference_conv_net(x_tf, keep_prob)
```

```
cost = loss_deep_conv_net(output, y_tf)
```

```
train_op = training(cost)
```

```
eval_op = evaluate(output, y_tf)
```

```
#####
```

```
## for metrics
```

```
y_p_metrics = tf.argmax(output, 1)
```

```
#####
```

```
# Initialize and run
```

```
#init = tf.global_variables_initializer()
```

```
init = tf.initialize_all_variables()
```

```
sess = tf.Session()
```

```
sess.run(init)
```

```
#####
```

```
# one-hot encoding
```

```
depth = 4
```

```
y_train_onehot = sess.run(tf.one_hot(y_train, depth))
```

```
y_test_onehot = sess.run(tf.one_hot(y_test, depth))
```

```
#####
```

```
## batch parameters
```

```
num_samples_train = len(y_train)
```

```
print num_samples_train
```

```
num_batches = int(num_samples_train/batch_size)
```

```
#####
```

```
#####
```

```
dropout2 = 1.0
```

```
#####
```

```
# MAIN_LOOP()
```

```
for i in range(n_epochs):
```

```
    for batch_n in range(num_batches):
```

```
        sta= batch_n*batch_size
```

```
        end= sta + batch_size
```

```
        sess.run( train_op , feed_dict={x_tf: x_train[sta:end,:], y_tf: y_train_onehot[sta:end, :], keep_prob: dropout })
```



```

Player
Terminal File Edit View Search Terminal Help
2 3 2 1 0 3 3 1 0 3 3 0 0 3 3 0 0 1 2 0 3 0 1 0 1 2 3 3 3 2 3 0 3 3 3 3 1
3 1 2 3 3 3 2 1 3 3 2 0 0 2 1 2 3 3 0 2 0 1 2 0 0 0 0 0 2 3 1 1 0 3 2 3 2
2 3 0 0 0 3 2 1 2 2 0 1 0 0 0 2 1 0 1 2 1 2 1 1 3 3 3 1 0 2 2 1 1 0 0 1 0
3 2 3 3 2 0 1 3 1 1 2 1 0 1 0 1 3 1 1 3 1 0 0 3 0 1 1 3 1 3 0 3 3 1 2 0 1
0 3 3 3 2 3 0 1 0 3 3 3 2 3 2 2 0 2 0 3 0 3 2 3 2 2 3 2 3 1 1 3 1 3 1 3 1
0 0 2 3 1 1 0 3 3 3 0 3 1 3 3 0 2 1 0 3 3 3 1 3 0 3 2 1 3 1 1 3 3 2 3 3 1
1 3 2 1 0 1 0 1 0 3 2 0 0 1 0 3 2 3 0 3 1 1 2 1 3 3 3 2 3 1 3 3 0 2 1 0 0
2 2 0 2 1 1 3 3 1 0 0 1 3 1 2 0 1 2 0 2 1 3 0 1 2 1 2 1 0 3 3 0 3 0 0 1 2
3 3 1 0 1 1 2 3 0 1 2 1 1 3 3 3 2 0 2 2 2 1 0 0 0 2 3 0 1 3 1 3 3 3 1 1 0
3 2 0 3 1 3]
[1 2 0 0 3 1 1 0 1 0 1 0 1 0 1 0 0 3 2 1 1 0 0 0 3 1 3 0 2 3 2 3 2 0 1 0 2
1 1 0 3 3 3 1 0 1 1 1 3 3 1 0 0 1 0 3 0 1 0 0 0 0 0 1 3 3 3 2 0 2 3 1 0 1 3
0 0 1 1 0 3 3 0 1 0 2 0 2 1 3 2 0 3 2 0 0 0 1 3 2 0 0 0 0 1 0 0 0 0 0 2 2 1 1
3 0 0 2 0 0 2 3 0 3 1 3 2 1 0 2 1 1 3 2 3 1 1 2 1 2 0 0 0 3 0 2 0 0 3 1 3
2 0 1 1 1 1 1 2 0 3 0 0 1 3 3 1 0 0 2 2 1 3 1 3 2 2 1 3 2 1 1 2 2 0 0 2 0
0 2 1 0 2 3 0 1 1 1 3 0 3 3 3 0 2 2 1 2 0 2 0 1 1 0 2 1 0 3 2 1 2 3 0 2 1
0 2 2 0 3 2 0 1 0 3 3 0 1 3 1 3 3 3 1 3 2 2 2 3 2 0 3 3 3 0 2 2 2 3 3 0
3 1 1 3 1 3 0 0 1 1 1 3 3 1 0 1 1 0 1 3 2 0 1 1 2 3 1 3 1 0 2 3 1 0 0 2
1 3 1 2 3 1 0 1 0 3 0 3 2 1 1 3 3 1 1 0 2 3 0 0 3 1 1 1 0 2 0 2 1 2 3 1 0
1 1 1 3 2 1 0 3 0 2 2 0 1 1 3 1 1 2 3 2 3 1 0 1 3 3 1 2 1 0 0 0 1 1 1 3 3
0 1 0 0 2 0 0 2 1 1 2 1 1 3 1 0 0 1 2 3 0 2 0 0 1 3 3 3 2 1 1 0 1 1 3 2
0 3 0 2 3 2 3 3 2 1 3 1 2 2 3 0 1 0 1 2 3 2 0 3 0 2 3 2 3 2 2 2 1 1 2 0
3 1 1 1 2 1 3 1 1 3 2 0 1 0 2 2 0 0 3 1 0 2 0 0 3 0 2 3 2 3 0 1 1 2 1 0 1
0 1 3 3 0 2 2 2 1 1 3 0 1 1 1 0 1 1 1 0 2 2 0 3 1 3 3 0 2 0 3 3 2 3 2 2 1 3
1 1 2 0 2 3 2 0 3 1 1 1 0 0 2 2 3 0 2 1 2 1 0 0 3 3 2 2 3 0 0 1 2 0 0 1
2 1 3 0 3 0 3 0 2 3 1 0 1 0 0 0 2 2 3 0 0 2 3 3 3 0 1 3 3 3 2 2 1 2 2 1 2
2 3 3 0 0 2 2 1 0 2 3 0 0 2 3 0 0 1 2 0 3 0 2 1 3 2 3 3 1 2 2 0 3 3 3 2 1
3 1 2 2 0 0 2 3 3 3 0 1 1 1 3 3 3 0 3 0 1 2 0 0 0 0 0 2 3 1 1 0 3 2 3 2
3 3 0 0 0 3 2 0 2 3 0 1 0 0 0 2 1 0 1 3 1 2 2 1 3 3 1 1 0 2 1 1 1 1 0 2 0
3 3 2 2 2 1 1 3 1 1 2 2 0 1 1 1 3 1 1 2 1 0 0 3 0 1 1 1 2 3 1 1 2 1 2 0 0
0 3 3 2 3 0 1 0 3 3 1 2 3 1 2 0 2 0 3 0 3 2 2 0 3 3 2 3 3 1 3 1 3 1 3 3
0 0 2 3 1 3 0 3 1 2 1 3 3 3 2 0 2 1 1 3 3 3 1 3 0 2 2 1 2 3 1 3 3 2 3 3 3
1 3 2 1 0 1 0 1 0 1 2 0 0 1 0 3 2 3 0 1 1 1 2 1 2 3 2 1 3 1 3 3 0 2 1 0 1
2 2 0 2 1 0 2 3 0 0 0 1 3 1 2 0 1 2 0 2 1 3 0 1 1 3 2 1 0 2 3 0 3 0 0 1 2
1 3 3 0 1 1 3 3 0 3 2 1 1 3 2 2 3 0 1 2 1 3 0 0 0 2 3 0 1 3 1 3 3 1 1 2 0
3 0 0 2 3 3]
Accuracy: 0.76
confusion matrix
[[225 16 4 3]
 [ 24 178 17 32]
 [ 0 16 125 58]
 [ 0 20 29 184]]
Predicted 0 1 2 3 All
True
0 225 16 4 3 248
1 24 178 17 32 251
2 0 16 125 58 199
3 0 20 29 184 233
All 249 230 175 277 931
Precision: 0.768
Recall: 0.765
F1-measure: 0.764
*****
terminate called after throwing an instance of 'std::bad_alloc'

```

Output:

RESULTS:

	PRECISION	RECALL	ACCURACY	F- MEASURE
COVOLUTION NEURAL NETWORK	0.768	0.764	0.76	0.764